

Realtime Events AI Triggers

From event streams → in-the-moment outcomes

A real-time, AI-personalized action layer for the streaming events you already collect.



Customers stream events. They can't reason over the last few minutes of them.

Real-time pipelines deliver events. Warehouse audiences segment customers over weeks. There's nothing in between that holds short-term memory of a session, watches for live patterns, and acts in seconds — so teams stitch it together with brittle scripts and miss the moment.

MARKETING

“

We see anonymous high-intent visitors leave every day. By the time tomorrow's report runs, they're gone.”

Cost

Missed conversions you can never re-attribute.

CUSTOMER SUCCESS

“

A whale browsed our pricing page yesterday. We found out at next week's QBR — too late to call.”

Cost

Deals sized in the room, not in real time.

ONBOARDING

“

Customers churn at the destination-connect step. Every error needs a different fix; we can't follow up on each one.”

Cost

Drop-off the moment users hit friction.

Six moments worth catching. One engine catches them all.



Capture anonymous high-intent

Visitor narrows filters, opens 3 listings, idles. Fire an in-app banner offering a tour or chat — before they leave the site, while they're still on it.

In-app capture moment

→ Marketing pain



Alert sales / CS on known engagement

Logged-in customer revisits pricing. Slack the right account owner with full context: name, ARR, prior interactions, the page they're dwelling on.

Real-time human handoff

→ Customer Success pain



Rescue stuck onboarding errors

Setup error fires. We send a personalized fix email referencing the exact error code, the destination type, and the customer's stack — not a generic FAQ.

Personalized per error

→ Onboarding pain



Re-engage stalled multi-step flows

Customer created a source but never connected a destination. After idle, an email referencing their progress and tech stack — not a generic nudge.

Mid-funnel context-aware

→ Onboarding pain



Synthetic events from event patterns

Emit a new derived event when a pattern matches — e.g., "high_intent_session" when 3 listings + a filter + 8s idle. Feeds back into the same SDK / warehouse.

Custom hooks on patterns

→ Platform extensibility



Temporal context API

Ask "what did this user do in the last 10 minutes?" — short-term memory exposed as a query, so agents and apps can reason over live behavior, not yesterday's batch.

Live session lookup

→ Agent / app context

+ All single-event "webhook on event X" patterns you'd otherwise wire by hand in Transformer — same engine, no new plumbing.

A 5-stage pipeline. Stateful. Plug-in.



✓ **12 composable predicates**

window.idle_seconds, traits.known,
has_event_name, AND/OR/NOT — fits any
pattern.

✓ **Templated AI actions**

Slack and email fill {{trait.first_name}} and
{{window.last_listing}} at fire time.

✓ **Pluggable destinations**

Slack + email today. Swap LLM_MODE=live,
plug your destination registry tomorrow.



LIVE DEMO → 3 concurrent visitors, two Slack pings, one personalized email, ~30 seconds

No one composes all six. We're the gap-filler.

Capability	Segment	Tealium	Customer.io	Braze	This
Single-event webhooks	●	●	●	●	●
Stateful windowing	×	◐	●	●	●
Idle / dwell predicates	×	×	◐	◐	●
Anonymous-intent action	×	◐	×	×	●
LLM-personalized content	×	×	×	×	●
Pluggable destinations	●	●	×	×	●

● Full ◐ Partial × Not supported

Bottom line: Other tools each solve one or two of these. None compose all six into a single substrate — which is why customers DIY this with cron jobs and Looker.

Hackathon code today. Production substrate tomorrow.

SHIPPED THIS WEEK

- ✓ Rules engine: 12 predicates, AND/OR/NOT, hot-reloadable from YAML
- ✓ Window manager: idle/dwell/counters, sharded, 15-min retention
- ✓ Cooldowns + dedup so the same realtor isn't paged twice
- ✓ Mock Activation API (real-shape compatible — single flag swap)
- ✓ Templated canned LLM responses + live local-agent fallback
- ✓ Slack webhook + email outbox + onboarding wizard + dashboard

PRODUCTION ROADMAP

- Multi-tenant + horizontal scaling (consumer per tenant)
- Live LLM mode — BYOK or platform-managed, with caching
- Real Activation API integration (already wire-compatible)
- ClickHouse for events archive + replay
- CEL or Drools for richer rule expressions (sandbox-safe)
- Pluggable destination registry: any webhook or in-app SDK



The AI-native automation substrate. Same engine, every event, every customer.